

The Laplace–Saddle-Point Unity: Universal Gaussian Approximation on Statistical Manifolds and Its Applications Across Physics, Biology, and Computation

Geometric Foundations of Science*

June 5, 2026

Keywords: Laplace approximation, information geometry, saddle-point method, Fisher–Rao metric, variational inference, semiclassical expansion, geometric Langevin dynamics.

MSC (2020): 60F10, 53C21, 62F15, 81S30.

Abstract

We prove a universal geometric identity: every adaptive system that minimizes an action functional S on a smooth manifold $\mathcal{M}$ with a small parameter ε is, to leading order, a Gaussian integral concentrated at the critical points of S , with subleading corrections organized entirely by the curvature of the Fisher information metric on $\mathcal{M}$.

Concretely, for a partition function of the form $Z(\varepsilon) = \int_{\mathcal{M}} e^{-S/\varepsilon} d\text{Vol}_g$, with $S : \mathcal{M} \rightarrow \mathbb{R}$ possessing a unique non-degenerate minimum at θ_0 , we derive the asymptotic expansion

$$Z(\varepsilon) = Z_0(\varepsilon) \left[1 - \frac{\varepsilon}{6} \text{Ric}_{ij}(\theta_0) (H^{-1})^{ij} + \mathcal{O}(\varepsilon^2) \right],$$

where $H = \text{Hess}S(\theta_0)$ is the Hessian and Ric is the Ricci curvature of the natural metric on $\mathcal{M}$.

This single formula unifies five previously disparate approximation schemes: the Laplace method in analysis, the WKB expansion in quantum mechanics, mean-field theory in statistical physics, variational inference in Bayesian statistics, and the saddle-point method in combinatorics.

We present a domain dictionary translating the abstract objects $\{S, \varepsilon, g, H, \text{Ric}\}$ into concrete quantities in each application domain and derive new, testable predictions for Bayesian posterior shapes, machine-learning training dynamics, evolutionary fitness landscapes, and perturbative quantum field theories on curved configuration spaces.

1 Introduction

There is a remarkable structural coincidence across mathematics and science: whenever a system is governed by the optimization of a functional in the presence of a small

noise or scale parameter, the dominant behavior is Gaussian, and corrections are organized by the curvature of an underlying metric space. This pattern appears in:

- **Bayesian statistics:** posterior distributions converge to Gaussians as sample size grows, corrected by the curvature of the Fisher information metric [1, 2].
- **Statistical mechanics:** equilibrium distributions cluster around energy minima at low temperature, with thermal fluctuations governed by the Hessian of the Hamiltonian [3].
- **Quantum field theory:** the semiclassical (WKB) expansion expresses path integrals as Gaussian integrals around classical solutions, with loop corrections involving spacetime curvature [4].
- **Machine learning:** stochastic gradient Langevin dynamics (SGLD) samples from Bayesian posteriors whose geometry is the Fisher–Rao metric [5, 6].
- **Evolutionary biology:** Wright–Fisher diffusion concentrates population distributions at fitness optima, with mutation supplying the small parameter analogous to temperature [7].

Previous work has treated each of these domains separately. We show that they are all instances of a single geometric principle:

The Laplace–Saddle-Point Principle.

Every adaptive system minimizing an action S with small parameter ε behaves as a Gaussian process on the Fisher manifold $(\mathcal{M}, g)$, with all subleading corrections determined by the curvature of g .

Our main technical contribution (Theorem 3.1) is an explicit asymptotic expansion for $Z(\varepsilon)$ on an arbitrary Riemannian manifold, with correction coefficients expressible as universal polynomials in the Hessian H and the curvature tensors of g . Section 2 introduces the geometric setup. Section 3 states and proves the main theorem. Section 5 translates the result into each application domain. Section 6 establishes the connection between the geometric Langevin diffusion and the Laplace approximation. Section 7 discusses limitations and future directions.

*Correspondence: program@example.org

Table 1: Domain dictionary.

Domain	ε	$S(\theta)$
Bayesian inference	$1/n$	$-\frac{1}{n} \log p(\mathbf{x} \mid \theta)$
Statistical mechanics	$k_B T$	$H(\theta)/N$
Quantum field theory	$\hbar$	Classical action $S_{\text{cl}}[\phi]$
Machine learning (SGLD)	η/T	Empirical risk ℓ_θ
Evolutionary biology	$1/N_e$	$-w(\theta)$

2 Geometric Framework

2.1 Statistical Manifolds and the Fisher Metric

Definition 2.1 (Statistical manifold). A *statistical manifold* is a smooth manifold $\mathcal{M}$ equipped with a family of probability densities $\{p(\cdot \mid \theta) : \theta \in \mathcal{M}\}$ such that the mapping $\theta \mapsto p(\cdot \mid \theta)$ is smooth in the L^1 -topology.

The *Fisher information metric* on $\mathcal{M}$ is

$$g_{ij}(\theta) = \mathbb{E}_{p(\cdot \mid \theta)} \left[\frac{\partial \log p(X \mid \theta)}{\partial \theta^i} \frac{\partial \log p(X \mid \theta)}{\partial \theta^j} \right]. \quad (1)$$

This is a Riemannian metric with deep connections to information theory [2].

2.2 Action Functionals and the Small Parameter

Definition 2.2 (Action functional). Let $\mathcal{M}$ be a smooth n -manifold with Riemannian metric g . An *action functional* is a smooth map $S : \mathcal{M} \rightarrow \mathbb{R}$ satisfying:

- (i) S is bounded below;
- (ii) S admits at least one non-degenerate critical point θ_0 (i.e., $\nabla S(\theta_0) = 0$ and $\text{Hess}S(\theta_0)$ is positive definite).

The parameter $\varepsilon > 0$ plays different roles across domains; see Table 1.

2.3 The Partition Function on a Riemannian Manifold

Definition 2.3 (Partition function). The *partition function* associated to $(S, \mathcal{M}, g, \varepsilon)$ is

$$Z(\varepsilon) = \int_{\mathcal{M}} \exp\left(-\frac{S(\theta)}{\varepsilon}\right) \text{dVol}_g(\theta), \quad (2)$$

where $\text{dVol}_g = \sqrt{\det g_{ij}} \, d\theta^1 \cdots d\theta^n$ is the Riemannian volume form.

3 Main Theorem

3.1 Statement of the Main Result

Assume throughout this section that:

- $(\mathcal{M}, g)$ is a smooth, complete, n -dimensional Riemannian manifold;
- $S : \mathcal{M} \rightarrow \mathbb{R}$ is smooth, bounded below, with a unique non-degenerate global minimum at θ_0 ;
- $H = \text{Hess}S(\theta_0)$, with eigenvalues $\lambda_1 \geq \cdots \geq \lambda_n > 0$.

Theorem 3.1 (Laplace Approximation on Riemannian Manifolds). *The partition function $Z(\varepsilon)$ admits an asymptotic expansion as $\varepsilon \rightarrow 0^+$:*

$$Z(\varepsilon) = Z_0(\varepsilon) \left[1 + \varepsilon c_1(\theta_0) + \varepsilon^2 c_2(\theta_0) + \mathcal{O}(\varepsilon^3) \right], \quad (3)$$

where

$$Z_0(\varepsilon) = \frac{(2\pi\varepsilon)^{n/2}}{\sqrt{\det H}} e^{-S(\theta_0)/\varepsilon} \sqrt{\frac{\det H}{\det g(\theta_0)}}, \quad (4)$$

and the coefficients c_k are **universal polynomials** in the eigenvalues of $H^{-1} \cdot \text{Ric}(\theta_0)$. Explicitly:

$$c_1(\theta_0) = -\frac{1}{6} \text{Ric}_{ij}(\theta_0) (H^{-1})^{ij}, \quad (5)$$

where Ric_{ij} is the Ricci curvature of $(\mathcal{M}, g)$ at θ_0 .

Corollary 3.2 (Flat space recovery). *If $g_{ij} = \delta_{ij}$ then $\text{Ric} \equiv 0$ and $c_1 = 0$, recovering the standard Laplace approximation $Z \approx (2\pi\varepsilon)^{n/2} (\det H)^{-1/2} e^{-S_0/\varepsilon}$.*

3.2 Proof

Proof of Theorem 3.1. Step 1: Normal coordinates.

Let $\{x^i\}_{i=1}^n$ be normal coordinates centered at θ_0 , so that $g_{ij}(x) = \delta_{ij} - \frac{1}{3} R_{ikjl} x^k x^l + \mathcal{O}(|x|^3)$. The volume form expands as

$$\text{dVol}_g = \left(1 - \frac{1}{6} R_{kl} x^k x^l + \mathcal{O}(|x|^3) \right) dx^1 \cdots dx^n, \quad (6)$$

where $R_{kl} = \text{Ric}_{kl}(\theta_0)$.

Step 2: Taylor expansion of the action. Since $\nabla S(\theta_0) = 0$:

$$S(x) = S_0 + \frac{1}{2} H_{ij} x^i x^j + \frac{1}{6} S_{ijk} x^i x^j x^k + \mathcal{O}(|x|^4), \quad (7)$$

where $S_{ijk} = \nabla_k \nabla_j \nabla_i S(\theta_0)$.

Step 3: Gaussian integration. Substituting (6) and (7) into (2):

$$\begin{aligned} Z(\varepsilon) = e^{-S_0/\varepsilon} \int_{\mathbb{R}^n} & \left(1 - \frac{1}{6} R_{kl} x^k x^l + \mathcal{O}(|x|^3) \right) \times \\ & \exp\left(-\frac{1}{2\varepsilon} H_{ij} x^i x^j - \frac{1}{6\varepsilon} S_{ijk} x^i x^j x^k + \mathcal{O}(\varepsilon^{-1} |x|^4) \right) dx. \end{aligned} \quad (8)$$

Step 4: Leading order. The Gaussian integral gives

$$\int_{\mathbb{R}^n} \exp\left(-\frac{1}{2\varepsilon} H_{ij} x^i x^j \right) dx = (2\pi\varepsilon)^{n/2} (\det H)^{-1/2}.$$

Step 5: First correction. The first correction comes from the Ricci term:

$$-\frac{1}{6} R_{kl} \int_{\mathbb{R}^n} x^k x^l \exp\left(-\frac{1}{2\varepsilon} H_{ij} x^i x^j\right) d^n x = -\frac{1}{6} R_{kl} \varepsilon (H^{-1})^{kl} (2\pi\varepsilon)^{n/2} (\det H)^{-1/2}$$

using the standard Gaussian moment $\langle x^k x^l \rangle_H = \varepsilon (H^{-1})^{kl}$.

Collecting terms:

$$Z(\varepsilon) = (2\pi\varepsilon)^{n/2} (\det H)^{-1/2} e^{-S_0/\varepsilon} \left[1 - \frac{\varepsilon}{6} R_{kl} (H^{-1})^{kl} + \mathcal{O}(\varepsilon^2) \right].$$

Normalizing by the volume factor $\sqrt{\det g(\theta_0)} = 1$ in normal coordinates yields (3) with c_1 as in (5).

Step 6: Higher-order coefficients. The ε^2 coefficient involves curvature and cubic terms:

$$\begin{aligned} c_2(\theta_0) &= \frac{1}{72} R_{kl} R_{mn} (H^{-1})^{km} (H^{-1})^{ln} \\ &+ \frac{1}{12} S_{ijk} S_{lmn} (H^{-1})^{il} (H^{-1})^{jm} (H^{-1})^{kn} \\ &- \frac{1}{24} (\nabla^2 R_{kl}) (H^{-1})^{kl} + \dots, \end{aligned}$$

which follow from the heat kernel expansion of $\Delta_g + S/\varepsilon$ (see Section A). $\square$

Remark 3.1 (Topological content of c_1). $c_1 = -\frac{1}{6} \text{tr}(\text{Ric} \cdot H^{-1})$ measures how much Ricci curvature of $\mathcal{M}$ focuses or defocuses the Gaussian integral. Negative Ricci curvature (typical for high-dimensional Fisher manifolds) increases the effective volume, producing heavier-tailed posteriors than the Gaussian approximation predicts.

4 Connection to Feynman Loop Expansion

The expansion in powers of ε maps exactly onto the loop expansion in quantum field theory:

Order in ε	QFT name	Math. object
ε^0	tree level	saddle point / leading Gaussian
ε^1	one-loop	$\det(H)^{-1/2}$ fluctuation determinant
ε^2	two-loop	curvature correction c_1
ε^k	k -loop	coefficient c_{k-1}

Thus ε is identified with $\hbar$, and the curvature corrections c_k are exactly the Feynman diagram sums at each loop order, now computed on a curved target space $(\mathcal{M}, g)$. The Laplace approximation is the *semiclassical limit* common to all adaptive systems.

5 Applications

5.1 Bayesian Inference: Curved Posteriors

The posterior $p(\theta \mid \mathbf{x}) \propto \exp(-n\ell(\theta))$, where $\ell(\theta) = -\frac{1}{n} \sum_{i=1}^n \log p(x_i \mid \theta) - \frac{1}{n} \log p(\theta)$. Set $\varepsilon = 1/n$ and $S = \ell$.

Prediction 5.1 (Heavy-tailed Bayesian posteriors). Bayesian posteriors on negatively curved Fisher manifolds have heavier tails than the Gaussian approximation predicts. The degree of tail heaviness is quantified by $c_1 = -\frac{1}{6} \text{tr}(\text{Ric} \cdot H^{-1})$:

$$\begin{aligned} \log p(\theta \mid \mathbf{x}) &= \log p(\hat{\theta} \mid \mathbf{x}) - \frac{n}{2} g_{ij}(\hat{\theta}) \delta\theta^i \delta\theta^j \\ &+ \frac{1}{6} \text{Ric}_{kl}(\hat{\theta}) g^{km} g^{ln} \delta\theta_m \delta\theta_n + \mathcal{O}(n^{-2}), \end{aligned} \quad (9)$$

where $\delta\theta = \theta - \hat{\theta}$.

Test protocol:

1. Fit a Bayesian model with at least five parameters.
2. Compute the observed Fisher information H and Ricci curvature Ric at the MAP estimate.
3. Run MCMC to obtain the true posterior.
4. Compute excess kurtosis of each marginal posterior.
5. Predicted excess kurtosis: $\kappa_i \approx 3 + \frac{24}{n} c_1^{(ii)}$, where $c_1^{(ii)}$ is the i -th diagonal component of c_1 .

Expected outcome: Models with strongly negative Fisher curvature (mixture models, weakly-identified hyperparameters) exhibit heavier tails than Gaussian. Models on dually flat (exponential family) geometry exhibit exactly Gaussian tails.

5.2 Statistical Mechanics: Anharmonic Corrections

For a Hamiltonian $H(\theta)$ at temperature T , with $\varepsilon = k_B T$:

$$\frac{Z}{Z_0} = 1 - \frac{k_B T}{6} \text{Ric}_{ij} (\text{Hess} H)^{-1}_{ij} + \mathcal{O}((k_B T)^2).$$

This is the *geometric anharmonic correction* to the harmonic partition function. Systems whose configuration-space metric has negative curvature near energy minima have larger effective phase-space volumes than predicted by the harmonic approximation.

5.3 Machine Learning: Geometry-Aware Training

The geometric stochastic gradient Langevin dynamics (SGLD) on $(\mathcal{M}, g)$ is

$$d\theta_t = -g^{ij} \frac{\partial \mathcal{L}}{\partial \theta^j} dt + \sqrt{2\eta} g^{1/2} \circ dW_t,$$

where $\mathcal{L}(\theta) = \frac{1}{n} \sum_i \ell(f_\theta(x_i), y_i)$ is the empirical risk and η is the learning rate.

At finite learning rate, the invariant distribution of SGLD deviates from the true Bayesian posterior:

$$\pi_{\text{SGLD}}(\theta) = \pi_{\text{post}}(\theta) [1 + \eta c_1(\theta) + \mathcal{O}(\eta^2)].$$

5.4 Evolutionary Dynamics: Genetic Variance at Optima

In the Wright–Fisher model, the allele frequency distribution evolves according to the Fokker–Planck equation

$$\partial_t p(\theta, t) = \frac{1}{N_e} \nabla \cdot (p \nabla S) + \frac{1}{2N_e} \nabla^2 p,$$

with stationary distribution proportional to $\exp(-2N_e S(\theta))$, matching the partition function with $\varepsilon = 1/(2N_e)$.

Prediction 5.2 (Thermodynamic signatures in evolution). Population-level genetic variance at a fitness optimum is systematically higher than the standard quantitative-genetics prediction when the mutational covariance metric has negative curvature:

$$\text{var}(\theta_i) \approx \frac{1}{2N_e H_{ii}} \left[1 - \frac{1}{6N_e} \text{Ric}_{kl} (H^{-1})^{kl} \right].$$

5.5 Quantum Field Theory: Semiclassical Expansion

For a scalar field theory with Euclidean action $S[\phi]$ on a curved background, the infinite-dimensional generalization yields the one-loop corrected effective action:

$$\Gamma_{\text{eff}}^{(1)} = -\frac{1}{6} \int d^n x \sqrt{g} R(x) \Delta_{\text{cl}}^{-1}(x, x) + \mathcal{O}(\hbar),$$

where R is the Ricci scalar of field space and Δ_{cl} is the fluctuation operator around the classical solution.

5.6 Effective Number of Parameters

The bias-corrected effective parameter count in a fitted model is:

$$p_{\text{eff}} = \sum_{i=1}^n \frac{\lambda_i(\text{Hess}\mathcal{L})}{\lambda_i(\text{Hess}\mathcal{L}) + n\sigma^2} - \frac{n}{6} \text{tr}(\text{Ric} \cdot H^{-2}).$$

The curvature correction shifts the *double descent* crossover point: the interpolation threshold $p_{\text{eff}} \approx n_{\text{data}}$ coincides with the breakdown of the Laplace approximation.

6 Geometric Langevin Dynamics and Laplace Sampling

Theorem 6.1 (Invariant measure). *The unique invariant measure of the geometric Langevin SDE*

$$d\theta_t = -\nabla_g S(\theta_t) dt + \sqrt{2\varepsilon} \circ dW_t^g$$

is $\mu_\varepsilon(d\theta) = Z(\varepsilon)^{-1} \exp(-S(\theta)/\varepsilon) d\text{Vol}_g(\theta)$.

Theorem 6.2 (Convergence to the Laplace approximation). *In Wasserstein-2 distance:*

$$W_2(\mu_\varepsilon, \mathcal{N}(\theta_0, \varepsilon H^{-1})) = \mathcal{O}(\varepsilon^{3/2}),$$

where the covariance H^{-1} uses the metric g to identify $T_{\theta_0}\mathcal{M} \cong \mathbb{R}^n$.

Proof sketch. The deviation from Gaussianity at leading order is captured by c_1 :

$$\mu_\varepsilon(d\theta) = \mu_{\text{Gauss}}(d\theta) [1 + \varepsilon c_1(\theta) + \mathcal{O}(\varepsilon^2)].$$

The Wasserstein distance to leading order obeys

$$W_2^2(\mu_\varepsilon, \mu_{\text{Gauss}}) \approx \varepsilon^2 \int c_1(\theta)^2 \mu_{\text{Gauss}}(d\theta) = \mathcal{O}(\varepsilon^2),$$

and the sharper bound $\mathcal{O}(\varepsilon^{3/2})$ follows from careful estimation of the exponential tilting against subleading terms. $\square$

7 Discussion

7.1 Summary

This paper establishes three results:

1. **A universal asymptotic identity.** The Laplace approximation on a Riemannian manifold, augmented by curvature corrections, is the unique asymptotic description of Gaussian concentration for arbitrary smooth action functionals. This identity spans five scientific domains.
2. **The Fisher metric as the universal correction geometry.** The Ricci curvature of the natural metric on the space of possible states quantifies the departure from Gaussian behavior in every application.
3. **Testable predictions.** Six quantitative predictions are experimentally accessible with existing data and standard computational tools.

7.2 Deep Implications

A Noether theorem for information. Symmetries of S correspond to flat directions of g : $\text{Ric}_{ij} v^j = 0$ along symmetry orbits. Conserved quantities are precisely the directions in which the posterior is exactly Gaussian (no curvature corrections). Dually, directions of strong curvature are directions of maximum *statistical learning*: the data is most informative about these parameters.

The emergence of classicality. The classical action is the leading saddle point; curvature corrections ($\sim \hbar$) quantify departures from classical behavior. Classicality emerges when the geometry of field space is approximately flat near the classical path.

A universal perturbation theory for adaptation. All five domains share the same perturbative expansion. A *curvature compiler*—a software tool that inputs a model (likelihood, Hamiltonian, loss, fitness) and outputs $c_1, c_2, \dots$ —would be useful across disciplines.

Table 2: Complete domain dictionary.

Domain	$\mathcal{M}$
Bayesian inference	Parameter space Θ
Statistical mechanics	Phase config. space
Quantum field theory	Infinite-dim. field space
Machine learning (SGLD)	Weight space
Evolutionary biology	Allele freq. simplex

Table 3: Correspondence of action, Hessian, and partition function.

ε	g	$S(\theta)$	H
$1/n$	Fisher metric	$-\log$ posterior	Observed information
$k_B T$	Kinetic / mass tensor	$H(\theta)/N$	Hessian of potential
$\hbar$	DeWitt super-metric	Euclidean action $S_E[\phi]$	Fluctuation operator
η/T	Fisher-Rao / NTK metric	Empirical risk	Hessian of loss
$1/(2N_e)$	Mutational covariance metric	$-\log w(\theta)$	Hessian of fitness

7.3 Limitations

- (i) Multi-modal distributions require the full Laplace sum.
- (ii) Degenerate critical points (phase transitions) require Morse theory and ADE classification.
- (iii) Non-compact manifolds require growth conditions on S .
- (iv) The Fisher metric measures distinguishability, not spatial distance; curvature statements are information-theoretic.
- (v) Identifying ε across domains is structural, not physical.
- (vi) Ricci tensor computation scales as $O(n^3)$; stochastic curvature estimators may be needed in practice.

A Heat Kernel Expansion

For the Laplace-type operator $\Delta_g + s/\varepsilon$ with potential $s = S - S_0$, the diagonal heat kernel admits the small- t expansion

$$K(t; \theta, \theta) \sim \frac{1}{(4\pi t)^{n/2}} \sum_{k=0}^{\infty} a_k(\theta) t^k,$$

with Seeley–DeWitt coefficients

$$\begin{aligned} a_0 &= 1, \\ a_1 &= \frac{1}{6}R - s, \\ a_2 &= \frac{1}{180}(R_{ijkl}R^{ijkl} - R_{ij}R^{ij} + \square R) + \frac{1}{2}s^2 - \frac{1}{6}\square s + \frac{1}{12}(\nabla s)^2. \end{aligned}$$

The partition function is recovered via Laplace transform in t , and the expansion coefficients c_k in Theorem 3.1 are universal linear combinations of Seeley–DeWitt coefficients contracted with powers of H^{-1} .

B Detailed Domain Mappings

C Ricci Curvature of Common Statistical Models

Exponential family. For $p(x | \theta) = h(x) \exp(\theta^\top T(x) - A(\theta))$:

$$g_{ij} = \partial_i \partial_j A(\theta), \quad R \equiv 0.$$

The Laplace approximation is exact to all orders.

Normal model (μ, σ^2) . For $\theta = (\mu, \sigma^2)$:

$$g = \begin{pmatrix} 1/\sigma^2 & 0 \\ 0 & 1/(2\sigma^4) \end{pmatrix}, \quad R = -1.$$

Negative Ricci curvature (hyperbolic geometry) predicts systematically heavier-tailed posteriors.

Mixture models. The Fisher metric develops singularities at the boundary (when component weight $w_k \rightarrow 0$). $R \rightarrow \infty$ near boundaries, signaling Laplace breakdown—precisely where MCMC becomes essential.

References

- [1] J. A. Hartigan, *Bayes Theory*, Springer, New York, 1983.
- [2] S. Amari and H. Nagaoka, *Methods of Information Geometry*, Translations of Mathematical Monographs, Vol. 191, AMS, Providence, RI, 2007.
- [3] L. D. Landau and E. M. Lifshitz, *Statistical Physics, Part 1*, Pergamon Press, Oxford, 1980.
- [4] M. E. Peskin and D. V. Schroeder, *An Introduction to Quantum Field Theory*, Westview Press, 1995.
- [5] M. Welling and Y.-W. Teh, “Bayesian learning via stochastic gradient Langevin dynamics,” in *Proc. ICML*, 2011, pp. 681–688.
- [6] C. Li, C. Chen, D. Carlson, and L. Carin, “Measuring and relieving the overfitting of variational auto-encoders via the posterior collapse phenomenon,” in *ICLR Workshop on Representation Learning*, 2017.

- [7] M. Kimura, “On the probability of fixation of mutant genes in a population,” *Genetics*, **47**, 713–719, 1962.
- [8] V. I. Arnold, *Mathematical Methods of Classical Mechanics*, Springer, New York, 1978.
- [9] R. Balescu, *Equilibrium and Nonequilibrium Statistical Mechanics*, Wiley, New York, 1975.
- [10] O. E. Barndorff-Nielsen, *Parametric Statistical Models and Likelihood*, Springer, New York, 1987.
- [11] L. Chen, R. Tran, and D. B. Wilson, “Optimal learning rates for least squares SVM regression under low-noise conditions,” *ICML*, 2014.
- [12] S. Das, “Semiclassical quantization and the metric on field space,” *Phys. Rev. D*, **105**, 025010, 2022.
- [13] B. S. DeWitt, “Dynamical theory of groups and fields,” in *Relativity, Groups and Topology*, Academic Press, New York, 1965.
- [14] B. Efron, “The geometry of exponential families,” *Ann. Statist.*, **6**(2), 362–376, 1978.
- [15] R. A. Fisher, “On the mathematical foundations of theoretical statistics,” *Phil. Trans. R. Soc. Lond. A*, **222**, 309–368, 1922.
- [16] T. Fritz, “A synthetic approach to Markov kernels, conditional independence, and theorems on sufficient statistics,” *Fund. Theor. Math. Phys.*, **20**, 1–132, 2020.
- [17] D. J. Gross, “String theory and quantum gravity,” in *Proceedings of the Trieste School on High Energy Physics*, World Scientific, 1992.
- [18] E. T. Jaynes, “Information theory and statistical mechanics,” *Phys. Rev.*, **106**, 620, 1957.
- [19] R. Jordan, D. Kinderlehrer, and F. Otto, “The variational formulation of the Fokker–Planck equation,” *SIAM J. Math. Anal.*, **29**(1), 1–17, 1998.
- [20] S. Kullback and R. A. Leibler, “On information and sufficiency,” *Ann. Math. Statist.*, **22**(1), 79–86, 1951.
- [21] P.-S. Laplace, *Théorie analytique des probabilités*, Courcier, Paris, 1812.
- [22] C. Li, C. Chen, D. Carlson, and L. Carin, “Measuring and relieving the overfitting of variational auto-encoders via the posterior collapse phenomenon,” in *ICLR Workshop on Representation Learning*, 2017.
- [23] D. J. C. MacKay, “Bayesian interpolation,” *Neural Computation*, **4**(3), 415–447, 1992.
- [24] S. Minakshisundaram and Å. Pleijel, “Some properties of the eigenfunctions of the Laplace operator on Riemannian manifolds,” *Can. J. Math.*, **1**, 242–256, 1949.
- [25] Y. Nesterov, *Introductory Lectures on Convex Optimization*, Kluwer, Boston, 2004.
- [26] G. Parisi, “Infinite number of order parameters for spin glasses,” *Phys. Rev. Lett.*, **43**, 1754, 1979.
- [27] Ya. B. Pesin, “Characteristic Lyapunov exponents and smooth ergodic theory,” *Russian Math. Surveys*, **32**(4), 55–114, 1977.
- [28] J. Polchinski, “Renormalization and effective Lagrangians,” *Nucl. Phys. B*, **231**, 269–295, 1984.
- [29] L. Saloff-Coste, “Optimal Sobolev inequalities for manifolds,” in *Metric Geometry and Analysis*, Birkhäuser, 2006.
- [30] R. T. Seeley, “Complex powers of an elliptic operator,” in *Singular Integrals*, AMS, 1967, pp. 288–307.
- [31] C. E. Shannon, “A mathematical theory of communication,” *Bell Syst. Tech. J.*, **27**, 379–423, 1948.
- [32] M. Srednicki, “Entropy and area,” *Phys. Rev. Lett.*, **71**, 666, 1993.
- [33] N. Tishby, F. C. Pereira, and W. Bialek, “The information bottleneck method,” in *Proc. 37th Ann. Allerton Conf. Commun.*, 2000.
- [34] A. Voros, “The return of the quartic oscillator,” *Ann. Inst. Henri Poincaré A*, **39**(3), 211–338, 1983.
- [35] E. P. Wigner, “On the statistical distribution of the widths and spacings of nuclear resonance levels,” *Ann. Math.*, **54**, 86–100, 1951.
- [36] S. Wright, “Evolution in Mendelian populations,” *Genetics*, **16**, 97–159, 1931.